

## DS 4300: Large Scale Information Storage and Retrieval Foundations

### Searching

- Searching is the most common operation performed by a database system.
- In SQL, **SELECT** is arguably the most versatile and complex statement.
- **Baseline for efficiency**: Linear Search (start at beginning and go element by element).
- **Record** = a collection of attribute values for a single entity instance (a row).
- **Collection** = a set of records of the same type (a table).
- **Search Key** = a value for an attribute used for searching.

### Lists of Records

- If each record takes  $x$  bytes,  $n$  records need  $n * x$  bytes.
- **Contiguously Allocated List**: all  $n * x$  bytes allocated in one chunk.
- **Linked List**: each record needs  $x$  bytes + space for pointers.

#### Arrays vs Linked Lists:

- Arrays: Faster random access, slower insertion (except at the end).
- Linked Lists: Faster insertion anywhere, but slower random access.

### Binary Search

- Input: sorted array + a target value.
- Output: index of target, or -1 if not found.
- Time complexity (worst case):  $O(\log n)$ .

**Linear Search** worst case:  $O(n)$ .

### Back to Database Searching

- Suppose data is stored on disk sorted by some primary key (like **id**), making searching on **id** fast.
- Searching by a different attribute (like **specialVal**) would require a linear scan unless we create an additional data structure (index).
- We can't store the table physically sorted by multiple attributes at once without duplicating data.

### Potential Data Structures for Indexing

- An **array** of (**searchKey**, **rowNumber**) tuples sorted by **searchKey** → supports binary search, but inserting is slow (because you must keep it sorted).

- A **linked list** of ( `searchKey`, `rowNumber` ) → supports faster insertion but linear search.
- A **binary search tree** → can provide relatively fast insertion and lookup.  
(A deeper discussion leads to balanced trees like *B-trees*, *B+ trees*, etc.)

## 03 - Moving Beyond the Relational Model

### DS 4300: Moving Beyond the Relational Model

#### Benefits of the Relational Model

- Mostly standard data model and query language.
- ACID compliance (Atomicity, Consistency, Isolation, Durability).
- Handles large amounts of structured data, with lots of tooling and expertise available.

#### Relational Database Performance

- Mechanisms to increase efficiency include:
  - Indexing
  - Direct control over storage
  - Column-oriented storage vs. row-oriented storage
  - Query optimization
  - Caching/prefetching
  - Materialized views
  - Stored procedures
  - Data replication and partitioning

#### Transaction Processing

- **Transaction** = sequence of CRUD operations executed as a single unit (all succeed or all fail).
- Ensures data integrity, concurrency control, error recovery, and simpler error handling.

#### ACID Properties

1. **Atomicity**: A transaction is all-or-nothing.
2. **Consistency**: A transaction takes the database from one consistent state to another, respecting constraints.
3. **Isolation**: Concurrent transactions do not interfere with each other (avoid dirty reads, non-repeatable reads, phantom reads).
4. **Durability**: Once committed, the transaction's changes are permanent, even if there's a system failure.

## Example: ACID Transaction (Money Transfer)

*(SQL pseudo-code showing **START TRANSACTION**, updates to two accounts, checking for negative balance, **ROLLBACK** or **COMMIT**, etc.)*

## Limits of Relational Databases

- Some applications do not need the full strength of ACID.
- Joins can be expensive for certain access patterns.
- A lot of modern data is semi-structured or unstructured (JSON, XML).
- Horizontal scaling can be challenging.
- Some apps require extremely high performance or real-time access.

## Scalability: Scale Up vs. Scale Out

- Traditional approach: scale vertically (bigger machines) until forced to scale out due to availability concerns.
- Horizontal scaling can be more complex but is often necessary for very large systems.

## Distributed Data and the CAP Theorem

- **Distributed system:** A collection of independent computers that appear as a single system to the user.
- **CAP Theorem:** A distributed data store cannot provide more than two of:
  - Consistency
  - Availability
  - Partition Tolerance
- In practice, systems must make trade-offs.

## 04-B+Tree Walkthrough

### B+ Tree Example

- Order  $m=4$ .
- Insert a series of keys, showing how the tree grows and splits.
- B+ trees keep data in leaf nodes, with interior nodes holding copies of keys for navigation.
- Balanced, ensuring  $O(\log n)$  lookup and insertion.

*(Slides walk through example insertions of 42, 21, 63, 89, 35, 10, 27, 96, 30, 37, showing how leaves split and eventually how the root splits.)*

## 04 - Data Replication

### Replicating Data

### Distributing Data - Benefits

1. Scalability / High throughput (handle larger volumes).
2. Fault Tolerance / High Availability (system can continue if nodes fail).
3. Latency (place data copies closer to global users).

### Distributed Data - Challenges

- Consistency: updates must propagate across the network.
- Increased application complexity: the application code may need to handle distribution.

### Replication

- Replicas contain copies of the same data.

### Common Replication Strategies

- **Single Leader:** One node (leader) handles writes; other nodes (followers) replicate.
- **Multiple Leaders:** Multiple nodes can accept writes, but conflict resolution is more complex.
- **Leaderless:** All nodes can accept writes, often employing quorum reads/writes and version vectors.

### Leader-Based Replication

- All writes go to the leader, which sends replication updates to followers.
- Clients can read from leader or followers.
- Used by many relational and NoSQL systems (MySQL, PostgreSQL, MongoDB replica sets, Kafka, etc.).

### How Replication Info Is Transmitted

- *Statement-based:* Forward actual SQL statements (risk with non-deterministic functions).
- *Write-ahead Log (WAL) shipping:* Byte-level log of storage engine changes.
- *Logical (row-based) log:* Higher-level description of row changes, more flexible.
- *Trigger-based:* Custom triggers log changes to special tables.

### Synchronous vs Asynchronous Replication

- **Synchronous:** Leader waits for a follower's acknowledgment, ensuring immediate consistency but possibly higher latency.
- **Asynchronous:** Leader does not wait. Better performance but leads to *eventual consistency* across replicas.

## Handling Leader Failure

- Need a failover mechanism: pick a new leader, ensure no conflicting leaders, handle partial replication (some writes might be lost).

## Replication Lag

- In asynchronous replication, followers may be behind the leader. This is the *inconsistency window*.

## Consistency Models with Replicas

- **Read-after-Write Consistency:** Ensure a user sees their own writes immediately (can force reads from leader).
- **Monotonic Read Consistency:** Prevent a user from reading older data after newer data in consecutive queries.
- **Consistent Prefix Reads:** Ensure that reads respect the write order (no out-of-order writes).

# 05 - NoSQL Intro + KV DBs

## NoSQL

- “Not Only SQL”: broad category of non-relational DBs.
- Often address large-scale data, flexible schemas, and distributed data with relaxed consistency.
- Rely on **BASE**: Basically Available, Soft State, Eventual Consistency.

## ACID vs. BASE

- Traditional RDBMS focus on strong consistency (ACID).
- Many NoSQL systems adopt BASE (eventual consistency).

## Categories of NoSQL

- Key-Value Stores
- Document Stores

- Wide-Column Stores
- Graph Databases
- Others (time-series, etc.)

## Key-Value Databases

- Simple: **key** = **value**.
- Fast, typically O(1) lookups.
- Great for caching, session storage, shopping cart data, etc.

### Redis

- In-memory data structure store.
- Supports various data types (strings, lists, sets, sorted sets, hashes, etc.).
- Very high performance (100k+ ops/sec).
- Optional persistence to disk (snapshots or append-only file).

## Redis Data Types (Overview)

- **Strings**: Basic text/binary up to 512 MB.
- **Lists**: Linked lists of strings.
- **Hashes**: Field-value pairs (like an object).
- **Sets**: Unordered unique strings, with set operations (union, intersect, diff).
- **Sorted Sets**: Like sets but with a score for ordering.
- **Geo**: Geospatial indexes.
- **JSON**: Can store and manipulate JSON data (via Redis modules).

## Common Use Cases for KV Stores

- Caching
- Session management
- Real-time analytics
- Feature store in ML
- Queues (producer/consumer patterns)

## 05b - Redis in Docker

### Setting Up Redis in Docker

1. In Docker Desktop, search for the “Redis” image and click **Run**.
2. Optionally expose port 6379 to the host.
3. In a production environment, secure Redis with a password and do not expose it directly to the internet.

4. Use DataGrip (or another client) to connect to Redis on `localhost:6379`.

## 06 - Redis + Python

### **redis-py**

- Standard Python client library for Redis.
- Install via `pip install redis`.
- Basic usage: `python`

CopyEdit

```
import redis
```

- `r = redis.Redis(host='localhost', port=6379, db=0, decode_responses=True)`
- 
- `decode_responses=True` ensures strings are returned (instead of bytes).

### **Redis Command Highlights (Strings)**

- `set(key, value)`, `get(key)`, `incr(key)`, `decr(key)`, etc.
- `mset(...)` to set multiple keys at once, `mget(...)` to retrieve multiple.

### **Redis Command Highlights (Lists)**

- `lpush(key, *values)`, `rpush(key, *values)`, `lpop(key)`, `rpop(key)`, etc.
- `lrange(key, start, stop)` to retrieve sublist.

### **Redis Command Highlights (Hashes)**

- `hset(key, field, value)` or `hset(key, mapping={...})`
- `hget(key, field)`, `hgetall(key)`, etc.

### **Pipelines**

- Group multiple commands in a pipeline to reduce network round trips: `python`

CopyEdit

- ```
pipe = r.pipeline()
```
- `pipe.set("seat:1", "#1")`
  - `pipe.set("seat:2", "#2")`
  - `results = pipe.execute()`
  -

## Redis in ML/DS

- Used for storing and quickly retrieving features, session data, or model outputs in real time.

## 07 - Document DBs and Mongo

### Document Databases

- Store data as JSON or similar (BSON).
- Flexible schema: each document can have different fields.
- Great for applications passing JSON/XML data.

### MongoDB

- Leading document store built on BSON (binary JSON).
- Basic hierarchy:
  - **Database**
    - **Collection**
      - **Documents**
- Flexible, easy horizontal scaling, powerful query language (`find`, `insert`, `update`, etc.), and indexing support.

### MongoDB Features

- Rich query support (CRUD).
- Indexing on document fields.
- Replication via replica sets.
- Horizontal sharding / partitioning.

### MongoDB Editions



- **MongoDB Atlas:** fully managed in the cloud.
- **MongoDB Enterprise:** subscription-based, self-managed.
- **MongoDB Community:** free to use, self-managed.

## Interacting with MongoDB

- `mongosh`: CLI shell.
- MongoDB Compass: official GUI.
- 3rd party tools: DataGrip, etc.
- Programmatically: PyMongo (Python), Mongoose (Node.js), etc.

## Example: Docker + MongoDB

- Pull official image, run container, map port 27017.
- Set environment variables for root username/password if needed.
- Use MongoDB Compass or another tool to connect.

## Example Queries in MongoDB Shell

- `db.movies.find(...)`
- `db.movies.find({year: 2010})`
- `db.movies.find({year: 2010, $or: [{ "awards.wins": {$gte: 5}}, {"genres": "Drama"}]})`
- `db.movies.countDocuments({...})`
- Projections: `db.movies.find(<filters>, {name: 1, _id: 0})`

# 08 - PyMongo

## Connecting with PyMongo

```
python
CopyEdit
from pymongo import MongoClient

client = MongoClient('mongodb://
user_name:pw@localhost:27017')
db = client['ds4300']
collection = db['myCollection']
```

## Basic Operations

- **Insert a document:** python

CopyEdit

- ```
post = {
```
- "author": "Mark",
  - "text": "MongoDB is Cool!",
  - "tags": ["mongodb", "python"]
  - }
  - post\_id = collection.insert\_one(post).inserted\_id
  - print(post\_id)
  -

- **Find documents:** python

CopyEdit

- ```
db.movies.find({"year": 2000})
```
- 

- **Counting documents:** python

CopyEdit

- ```
db.collection.count_documents({})
```
- 

- *Etc.* for updates, deletes, indexing, etc.

## 09 - Introduction to Graph Data Model

**Title:** DS 4300 – Introduction to the Graph Data Model

**Key Points:**

### 1. Graph Databases

- Data model based on nodes (vertices) and edges (relationships).
- Nodes and edges each can contain properties (like key-value pairs).
- Edges connect nodes; edges are often directed and can have labels.

- Graph queries often involve traversals (finding paths, shortest paths, etc.).
- 2. **Where Graphs Show Up**
  - **Social networks** (modeling people and relationships).
  - **The Web** (pages as nodes, hyperlinks as edges).
  - **Chemical and biological data** (genes, proteins, chemical compounds, etc.).
- 3. **Basics of Graphs**
  - **Labeled Property Graph:** nodes have labels (grouping them), properties on nodes and edges.
  - Nodes can exist without relationships; edges must connect two nodes.
  - **Paths:** ordered sequence of connected nodes/edges with no repeated nodes or edges.
- 4. **Flavors of Graphs**
  - **Connected vs. Disconnected.**
  - **Weighted vs. Unweighted** (edges may have weights).
  - **Directed vs. Undirected** (edges may be directional).
  - **Acyclic vs. Cyclic.**
  - **Sparse vs. Dense.**
- 5. **Types of Graph Algorithms**
  - **Pathfinding** (shortest path, BFS, DFS, spanning tree, max flow, etc.).
  - **Centrality:** determining which nodes are most “important” (e.g., social media influencers).
  - **Community Detection:** clustering or grouping nodes (detecting subgraphs or communities).
- 6. **Famous Graph Algorithms**
  - **Dijkstra’s:** single-source shortest path with positive edge weights.
  - **A\*:** uses a heuristic to guide pathfinding.
  - **PageRank:** ranks nodes by importance based on incoming links from other “important” nodes.
- 7. **Neo4j**
  - A popular graph database that supports transactional and analytical workloads.
  - Similar alternatives include Microsoft Cosmos DB (graph features) and Amazon Neptune.

## 10 - Neo4j

**Title:** DS 4300 – Neo4j

**Key Points:**

### 1. Neo4j

- Graph DB system for transactional and analytical graph-based data.
- Schema-optional (can impose a schema if desired).
- ACID-compliant. Supports indexing, distribution, and sharding.

## 2. **Cypher Query Language**

- Introduced by Neo4j in 2011.
- Pattern-based syntax: `( nodes ) - [ :REL ] -> ( otherNodes )`.
- A visual, declarative query style for matching relationships in a graph.

## 3. **APOC Plugin**

- “Awesome Procedures on Cypher”: library with hundreds of extra procedures/functions.
- Extends functionality for data export/import, graph algorithms, integration tasks.

## 4. **Graph Data Science Plugin**

- Implements common graph algorithms (PageRank, community detection, etc.) efficiently within Neo4j.

## 5. **Docker Compose for Neo4j**

- Example `docker-compose.yml` snippet:
  - Uses `image: neo4j:latest` and sets ports `7474` (HTTP browser UI) and `7687` (Bolt protocol).
  - Exposes environment variables for password, enabling APOC, GDS plugins, etc.
  - Maps volumes for data, logs, import, and plugins.

## 6. **Neo4j Browser**

- Web UI on `localhost:7474` once the container is up.
- Can execute Cypher commands directly.

## 7. **Basic Cypher Examples**

- `CREATE (:User {name: "Alice", birthPlace: "Paris"})`

- Matching a node: `c`

[CopyEdit](#)

```
MATCH (u:User {birthPlace: "London"})
```

- `RETURN u.name, u.birthPlace`

- 

- Relationship creation: `cypher`

[CopyEdit](#)

- ```
MATCH (alice:User {name:"Alice"}), (bob:User {name:"Bob"})
```
- ```
CREATE (alice)-[:KNOWS {since: "2022-12-01"}]->(bob)
```
- 

## 8. Importing CSV

- Place CSV file in the `import` directory used by the Docker container.
- Example: `cypher`  

CopyEdit

- ```
LOAD CSV WITH HEADERS
```
- ```
FROM 'file:///netflix_titles.csv' AS line
```
  - ```
CREATE (:Movie { id: line.show_id, title: line.title, releaseYear: line.release_year })
```
  -
- Using `split(...)` + `UNWIND` to handle multiple directors in one record, possibly `MERGE` to avoid duplicates.

## 11 - AWS Intro

**Title:** DS 4300 – AWS Introduction

**Key Points:**

### 1. Amazon Web Services (AWS)

- A leading public-cloud platform with 200+ services.
- Global infrastructure of “Regions” and “Availability Zones.”
- Pay-as-you-use model.

### 2. Historical Context

- Launched in 2006 with S3 (object storage) and EC2 (virtual servers).
- Grew to include RDS, DynamoDB, CloudFront, and more.

### 3. AWS Service Categories

- **Compute** (e.g., EC2, Lambda, ECS, EKS).

- **Storage** (e.g., S3, EFS, EBS).
- **Databases** (Relational: RDS/Aurora; NoSQL: DynamoDB, DocumentDB, etc.).
- **Analytics** (EMR, Athena, Glue, Redshift, Kinesis).
- **Machine Learning** (SageMaker, Comprehend, Rekognition, etc.).
- 4. **Cloud Models (IaaS, PaaS, SaaS)**
  - **IaaS**: Infrastructure as a Service (e.g., EC2).
  - **PaaS**: Platform as a Service (removes infrastructure management).
  - **SaaS**: Ready-to-use software hosted by the vendor (e.g., Salesforce).
- 5. **Shared Responsibility Model**
  - **AWS**: Security “OF” the cloud (data centers, hypervisors, etc.).
  - **Customer**: Security “IN” the cloud (own data, IAM, encryption, OS patches if self-managed).
- 6. **AWS Global Infrastructure**
  - **Regions**: distinct geographic areas (e.g., `us-east-1`).
  - **Availability Zones**: isolated data centers in each region.
  - **Edge Locations**: content delivery endpoints (CDN, caching).
- 7. **Key Services**
  - **Compute**: EC2 (VMs), ECS/EKS (containers), Lambda (serverless).
  - **Storage**: S3 (object store), EBS (block store), EFS (shared file system).
  - **Databases**: RDS, DynamoDB, DocumentDB, Neptune, etc.
  - **Analytics**: Athena, EMR, Glue, Redshift.
  - **ML**: SageMaker, pre-trained AI services.
- 8. **AWS Free Tier**
  - Offers limited free usage for new accounts (12 months for select services).
  - Must watch for usage exceeding free tier to avoid unexpected charges.

## 12 - EC2 & Lambda

**Title:** DS 4300 – Amazon EC2 & Lambda

**Key Points:**

### Amazon EC2

1. **EC2 Overview**
  - “Elastic Cloud Compute”: scalable virtual machines in the cloud.
  - Many instance types with different CPU, memory, GPU, etc.
2. **EC2 Features**
  - **Elasticity**: programmatically scale up/down.
  - **AMI**: Amazon Machine Image for pre-configured OS or custom images.

- Integrates easily with other AWS services (S3, RDS, etc.).
- 3. **EC2 Lifecycle**
  - *Launch* (create a new instance).
  - *Start/Stop* (pause usage if instance is EBS-backed).
  - *Terminate* (delete the instance).
  - *Reboot* (restart, preserving data on root volume).
- 4. **EC2 Storage Options**
  - *Instance Store*: ephemeral storage tied to the instance's lifecycle.
  - *EBS* (Elastic Block Store): persistent block-level storage volumes.
  - *EFS* (Elastic File System): managed NFS-like file system for multiple instances.
  - *S3* for object storage or backups.
- 5. **EC2 Use Cases**
  - Web hosting, data processing, ML training, disaster recovery, etc.
- 6. **Example Setup**
  - Launching an EC2 instance with Ubuntu, installing software (conda, Streamlit, etc.), opening security group ports.

## **AWS Lambda**

1. **Lambda Overview**
  - Serverless compute: run code in response to events (HTTP requests, S3 file upload triggers, etc.).
  - Pay only for execution time, not for idle time.
2. **Lambda Features**
  - Event-driven (integrates with S3, SNS, API Gateway, etc.).
  - Supports multiple runtimes (Python, Node, Java, etc.).
  - Automatically scales.
3. **Workflow**
  - Upload code via AWS Console or CLI, configure triggers/events.
  - Lambda runs code on-demand.
4. **Example**
  - Creating a Lambda function in the AWS Console, writing Python code (handler function), configuring test events or real event triggers.

## **12.6. B-Trees — CS3 Data Structures & Algorithms**

**Source:** Summarized from a tutorial on B-trees and B+ trees.

**Key Points:**

## 1. B-Trees

- Proposed by Bayer & McCreight (1972) for disk-based search trees.
- Shallow (height-balanced) with potentially high branching factor, minimizing disk I/O.
- Each node can have multiple children; typically matches disk block size.

## 2. Properties of a B-tree of order m:

- The root has at least two children if it's not a leaf.
- Each internal (non-root) node has between  $\lceil m/2 \rceil$  and m children.
- All leaves are at the same level (height-balanced).
- Each node typically holds key-value pairs up to a certain maximum, ensuring each node is at least half-full.

## 3. Operations:

- *Search*: similar to searching in a 2-3 tree, do a binary search in the current node's keys, descend the appropriate child.
- *Insertion*: find the correct leaf; if the leaf is full, split it and promote the middle key to the parent, possibly recursively splitting upward.
- *Deletion*: if removing from a node leaves it underfull ( $< \text{half capacity}$ ), attempt to rebalance by borrowing from siblings or merging nodes.

## 4. B+ Trees

- A variant of B-trees (often used in databases).
- **Difference**: all actual records (data) are stored in leaf nodes; internal nodes store only keys for navigation.
- Leaves are often linked for range scans (doubly linked list).
- Also must remain half-full.
- Searching continues down to leaf level, even if a matching key is found in an internal node (since that internal node is just an index).

## 5. B+ Tree Insertion/Deletion

- Similar to B-tree logic but with the rule that data resides only in leaves, internal nodes hold guide keys.
- Leaves also have a linked-list chain for sequential access.
- Balanced structure ensures  $O(\log n)$  search, insertion, and deletion times.

# AVL-Tree-Rotations.pdf

This file is a **PNG image** illustrating **AVL tree rotations**. The essential knowledge (textual) about AVL rotations includes:

- **AVL Trees**: self-balancing binary search trees where the height of the left and right subtrees of any node differ by at most 1.
- **Rotations** fix imbalances. Key rotation patterns:
  - **LL rotation** (Right rotation).



- **RR rotation** (Left rotation).
- **LR rotation** (Left-Right).
- **RL rotation** (Right-Left).
- Each rotation re-links a small set of pointers in constant time, preserving the BST property while reducing subtree height.

*(The PNG visually shows how subtrees are reattached during rotations.)*

## B-trees.pdf

**Title:** “B-trees”

**Key Points:**

1. **Context:** B-trees store multiple elements in each node, exploiting locality (useful for caching or disk storage).
2. **Order m:** Non-leaf nodes can have up to m children, leaves hold the actual data.  
Invariants:
  - All leaves are at the same level (height-balance).
  - Each node is at least half-full (except possibly root).
  - The number of keys in an internal node is one less than the number of children.
3. **Operations:**
  - *Lookup:* find correct child pointer by comparing the search key to node keys, continuing until leaf is reached.
  - *Insertion:* insert into a leaf. If leaf is full, split, push up the middle key to the parent. Possibly recurse upward if the parent is full, etc.
  - *Deletion:* if a leaf goes below half-full, borrow from siblings or merge with a sibling; can cascade upward if the parent likewise becomes underfull.
4. **Benefits:**
  - Minimizes disk I/O by keeping branching factor large.
  - Balanced:  $O(\log n)$  worst-case height and search times.

## C12-bst.pdf

**Title:** Chapter 12 – Binary Search Trees

**Key Points:**

1. **BST-Property**
  - For any node x, every key in its left subtree  $< \text{key}[x]$ , and every key in its right subtree  $> \text{key}[x]$ .
2. **BST Operations**
  - **Traversal:** inorder, preorder, postorder.
    - Inorder of a BST yields sorted keys.

- **Search:** start from root, compare target key to current node's key, go left or right accordingly until found or nil.
- **Minimum/Maximum:** leftmost node is min, rightmost node is max.
- **Successor/Predecessor:** next larger (or smaller) in sorted order, found by a combination of descending right-then-left or parent references.
- **Insertion:** find the correct leaf position; place the new node there.
- **Deletion:**
  - No children: remove the node outright.
  - One child: splice out the node, promote its single child.
  - Two children: find the node's successor (or predecessor), copy that key/data into the node to be deleted, then remove the successor from its location (which is a simpler case, as successor has  $\leq 1$  child).

### 3. Complexity

- If the tree height is **h**, basic operations are  $O(h)$ . Balanced BSTs have  $h = O(\log n)$ , but worst-case (unbalanced) can be  $O(n)$ .

## ICS 46 Spring 2022, Notes and Examples: AVL Trees

Title: ICS 46 – AVL Trees

Key Points:

### 1. Binary Search Tree Balancing

- BST performance depends on shape (height). A “degenerate” BST can degrade operations to  $O(n)$ .

### 2. AVL Trees

- **Definition:** A BST where for every node, the heights of the left and right subtrees differ by at most 1.
- Ensures worst-case height is  $O(\log n)$ .

### 3. Rotation Operations

- **LL** (single right rotation), **RR** (single left rotation).
- **LR** (left then right rotation), **RL** (right then left).
- Rotations rebalance the tree in constant time by adjusting only a few pointers.

### 4. Insertion in an AVL Tree

- Insert as in a normal BST.
- Then “walk back up” to the root, checking for violations of the AVL balance property.
- If a node is imbalanced, apply one of LL, RR, LR, or RL rotations depending on the direction of insertion.

### 5. Deletion in an AVL Tree

- Delete as in a normal BST.

- Then walk back up, re-check balance factors, applying rotations where needed.
- Maintaining the node's heights in each node helps check balance quickly.

6. **Why a “near-balance” criterion?**

- Perfect or complete balance is expensive to maintain for every insertion.
- AVL's local rebalancing ensures  $O(\log n)$  height, with  $O(\log n)$  insertion/deletion overhead.